

Introduction to  
**Information Retrieval**

Boolean Retrieval

# Terminology

- In the context of a **user** interacting with an IR system
  - **Document**: unit of retrieval
  - Each document has a **Doc Id**
  - **Corpus**: collection of documents
  - User has **information need**
  - User inputs a **query** to system
  - Term: a unit of information (e.g., a word/phrase)
    - **Relevance** of documents to query/info need
- **Ad hoc retrieval task**

# For most of this lecture

- Corpus: collection of plays of Shakespeare
- Document: an individual play
- Query: a Boolean expression having terms connected with Boolean operators (AND, OR, NOT)

# Unstructured data in 1620

- Which plays of Shakespeare contain the words ***Brutus AND Caesar*** but ***NOT Calpurnia***?
- One could grep all of Shakespeare's plays for ***Brutus*** and ***Caesar***, then strip out lines containing ***Calpurnia***?
- Why is that not the answer?
  - Slow (for large corpora)
  - **NOT *Calpurnia*** is non-trivial
  - Other operations (e.g., find the word ***Romans*** near ***countrymen***) not feasible
  - Ranked retrieval (best documents to return)
    - Later lectures

# Term-document incidence matrices

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

*Brutus AND Caesar BUT NOT Calpurnia*

1 if play contains word, 0 otherwise

# Incidence vectors

- So we have a 0/1 vector for each term.
- To answer query: take the vectors for ***Brutus***, ***Caesar*** and ***Calpurnia*** (complemented) → bitwise AND.
  - 110100 AND
  - 110111 AND
  - 101111 =
  - **100100**

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
<b>Antony</b>	1	1	0	0	0	1
<b>Brutus</b>	1	1	0	1	0	0
<b>Caesar</b>	1	1	0	1	1	1
<b>Calpurnia</b>	0	1	0	0	0	0
<b>Cleopatra</b>	1	0	0	0	0	0
<b>mercy</b>	1	0	1	1	1	1
<b>worser</b>	1	0	1	1	1	0

# Answers to query

- **Antony and Cleopatra, Act III, Scene ii**

*Agrippa* [Aside to DOMITIUS ENOBARBUS]: Why, Enobarbus,  
When Antony found Julius **Caesar** dead,  
He cried almost to roaring; and he wept  
When at Philippi he found **Brutus** slain.

- **Hamlet, Act III, Scene ii**

*Lord Polonius*: I did enact Julius **Caesar** I was killed i' the  
Capitol; **Brutus** killed me.

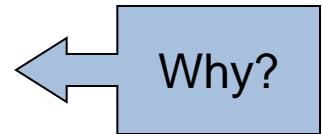


# Bigger collections

- Consider  $N = 1$  million documents, each with about 1000 words.
- Avg 6 bytes/word including spaces/punctuation
  - 6GB of data in the documents.
- Say there are  $M = 500K$  *distinct* terms among these.

# Can't build the matrix

- $500K \times 1M$  matrix has half-a-trillion 0's and 1's.
- But it has no more than one billion 1's.
  - matrix is extremely sparse.
- What's a better representation?
  - We only record the 1 positions.



# Introduction to **Information Retrieval**

The Inverted Index

The key data structure underlying  
modern IR

# Inverted index

- For each term  $t$ , we must store a list of all documents that contain  $t$ .
  - Identify each doc by a **docID**, a document serial number
- Can we use fixed-size arrays for this?

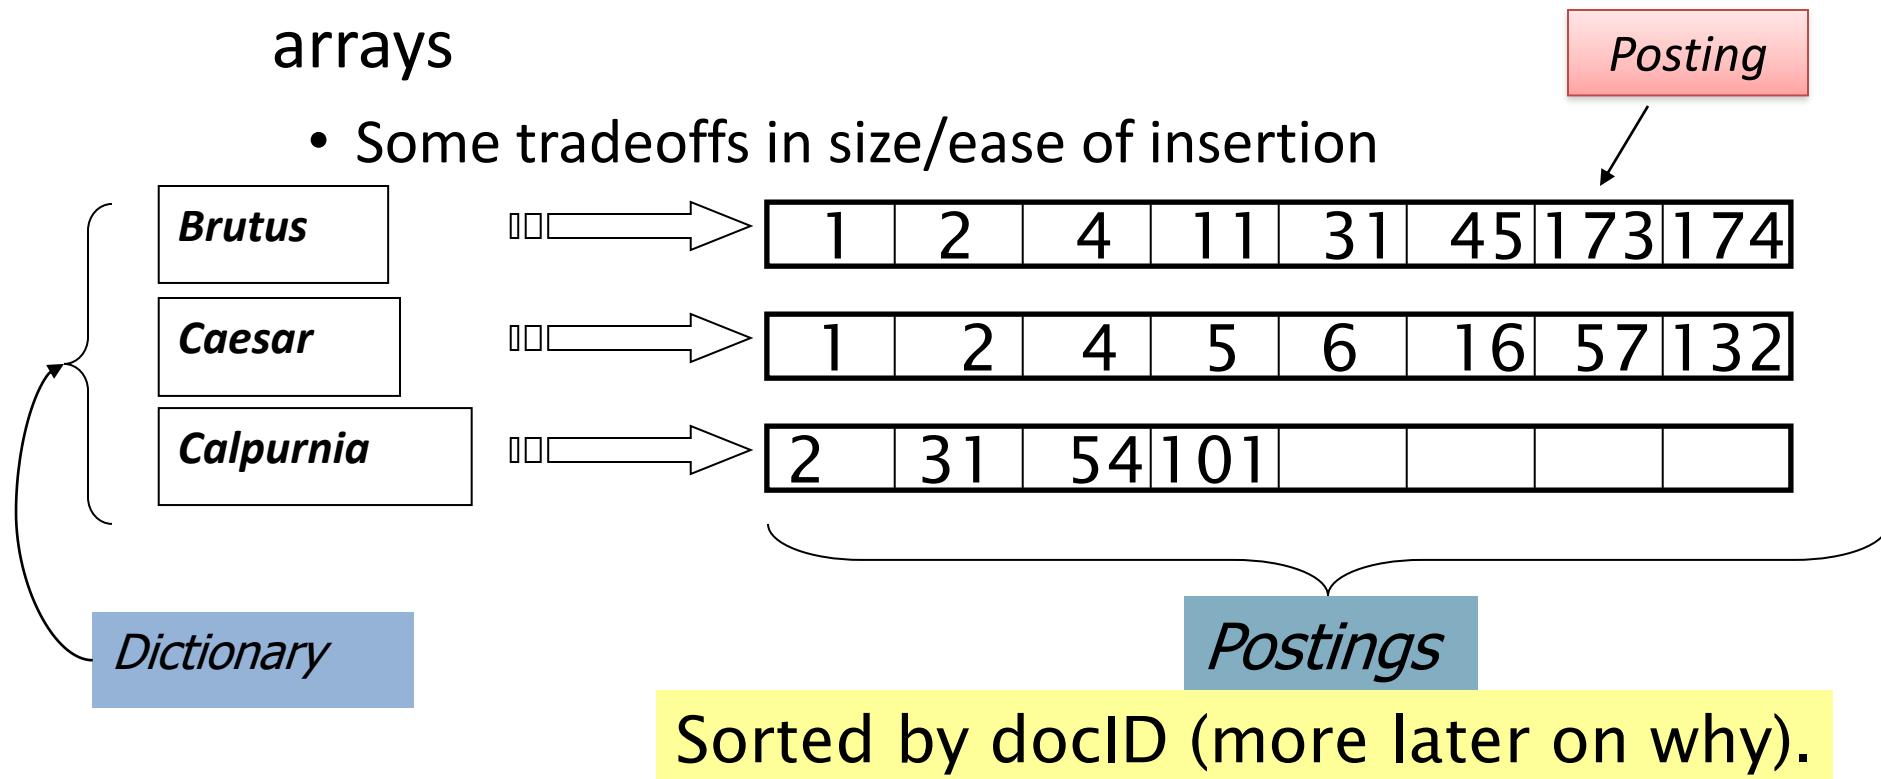
<b>Brutus</b>	⇒	1 2 4 11 31 45 173 174
<b>Caesar</b>	⇒	1 2 4 5 6 16 57 132
<b>Calpurnia</b>	⇒	2 31 54 101

What happens if the word **Caesar** is added to document 14?

# Inverted index

- We need variable-size **postings lists**
  - On disk, a continuous run of postings is normal and best
  - In memory, can use linked lists or variable length arrays

- Some tradeoffs in size/ease of insertion



# Inverted index construction

Documents to be indexed



Friends, Romans, countrymen.

⋮

Tokenizer

Token stream

Friends

Romans

Countrymen

Linguistic modules

Modified tokens

friend

roman

countryman

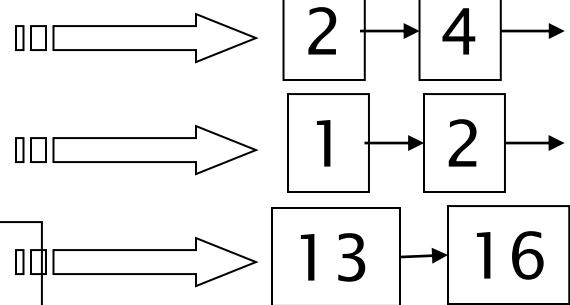
Indexer

Inverted index

*friend*

*roman*

*countryman*



# Initial stages of text processing

- Tokenization
  - Cut character sequence into word tokens
    - Deal with “*John’s*”, *a state-of-the-art solution*
- Normalization
  - Map text and query term to same form
    - You want *U.S.A.* and *USA* to match
- Stemming
  - We may wish different forms of a root to match
    - *authorize, authorization*
- Stop words
  - We may omit very common words (or not)
    - *the, a, to, of*

# Indexer steps: Token sequence

- Sequence of (Modified token, Document ID) pairs.

Doc 1

I did enact Julius  
Caesar I was killed  
i' the Capitol;  
Brutus killed me.

Doc 2

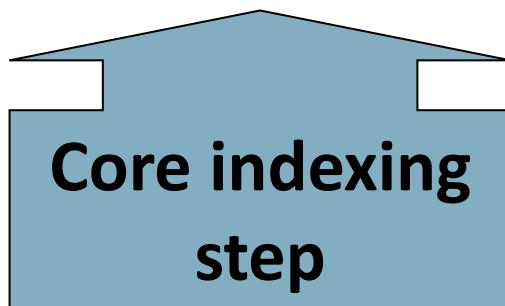
So let it be with  
Caesar. The noble  
Brutus hath told you  
Caesar was ambitious



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

# Indexer steps: Sort

- Sort by terms
  - And then docID



Term	docID
I	1
did	1
enact	1
julius	1
caesar	1
I	1
was	1
killed	1
i'	1
the	1
capitol	1
brutus	1
killed	1
me	1
so	2
let	2
it	2
be	2
with	2
caesar	2
the	2
noble	2
brutus	2
hath	2
told	2
you	2
caesar	2
was	2
ambitious	2

Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
I	1
I	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2

# Indexer steps: Dictionary & Postings

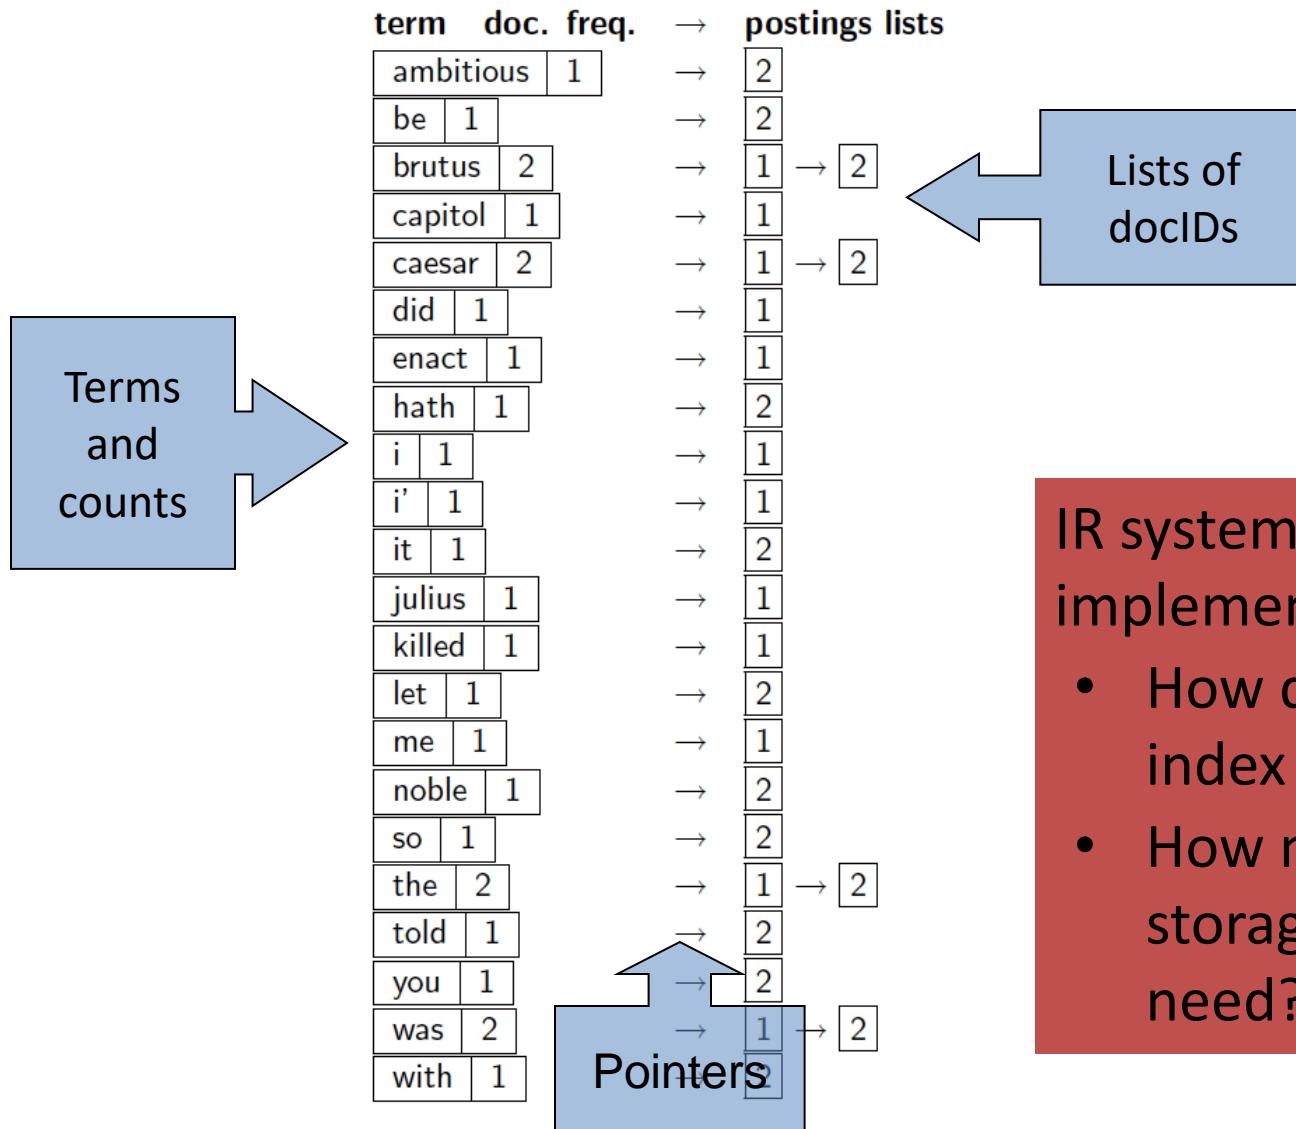
- Multiple term entries in a single document are merged.
- Split into Dictionary and Postings
- Document frequency information is added to dictionary.

Term	docID
ambitious	2
be	2
brutus	1
brutus	2
capitol	1
caesar	1
caesar	2
caesar	2
did	1
enact	1
hath	1
i	1
i	1
i'	1
it	2
julius	1
killed	1
killed	1
let	2
me	1
noble	2
so	2
the	1
the	2
told	2
you	2
was	1
was	2
with	2



term	doc.	freq.	→	postings lists
ambitious		1	→	2
be	1		→	2
brutus		2	→	1 → 2
capitol		1	→	1
caesar		2	→	1 → 2
did	1		→	1
enact		1	→	1
hath		1	→	2
i	1		→	1
i'	1		→	1
it	1		→	2
julius		1	→	1
killed		1	→	1
let	1		→	2
me		1	→	1
noble		1	→	2
so		1	→	2
the		2	→	1 → 2
told		1	→	2
you		1	→	2
was		2	→	1 → 2
with		1	→	2

# Where do we pay in storage?



## IR system implementation

- How do we index efficiently?
- How much storage do we need?

# Practical considerations

- For a practical IR system handling a huge corpus
  - The dictionary will be stored in the memory
  - Postings lists will be stored on disk
  - Ideally, retrieve (from disk) only those postings lists that are needed to answer a query

Introduction to  
**Information Retrieval**

Query processing with an inverted index

# The index we just built

- How do we process a query?
  - Later - what kinds of queries can we process?



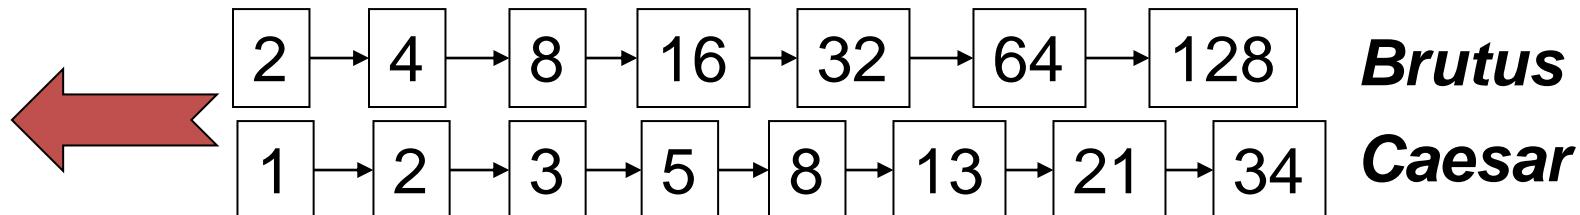
***Brutus AND Caesar***

# Query processing: AND

- Consider processing the query:

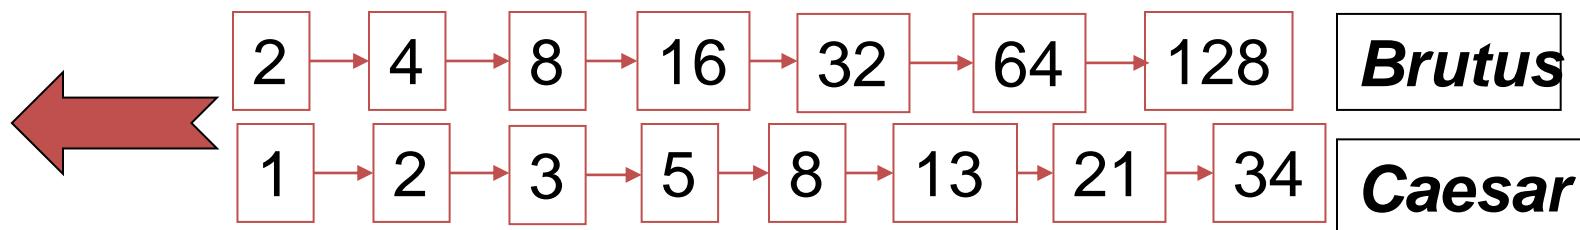
***Brutus AND Caesar***

- Locate ***Brutus*** in the Dictionary;
  - Retrieve its postings.
- Locate ***Caesar*** in the Dictionary;
  - Retrieve its postings.
- “Merge” the two postings (intersect the document sets):



# The merge

- Walk through the two postings simultaneously, in time linear in the total number of postings entries



If the list lengths are  $x$  and  $y$ , the merge takes  $O(x+y)$  operations.

Crucial: postings sorted by docID.

# Intersecting two postings lists (a “merge” algorithm)

INTERSECT( $p_1, p_2$ )

```
1  answer ← ⟨ ⟩  
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$   
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$   
4      then ADD(answer,  $\text{docID}(p_1)$ )  
5           $p_1 \leftarrow \text{next}(p_1)$   
6           $p_2 \leftarrow \text{next}(p_2)$   
7      else if  $\text{docID}(p_1) < \text{docID}(p_2)$   
8          then  $p_1 \leftarrow \text{next}(p_1)$   
9          else  $p_2 \leftarrow \text{next}(p_2)$   
10 return answer
```

# Boolean queries: Exact match

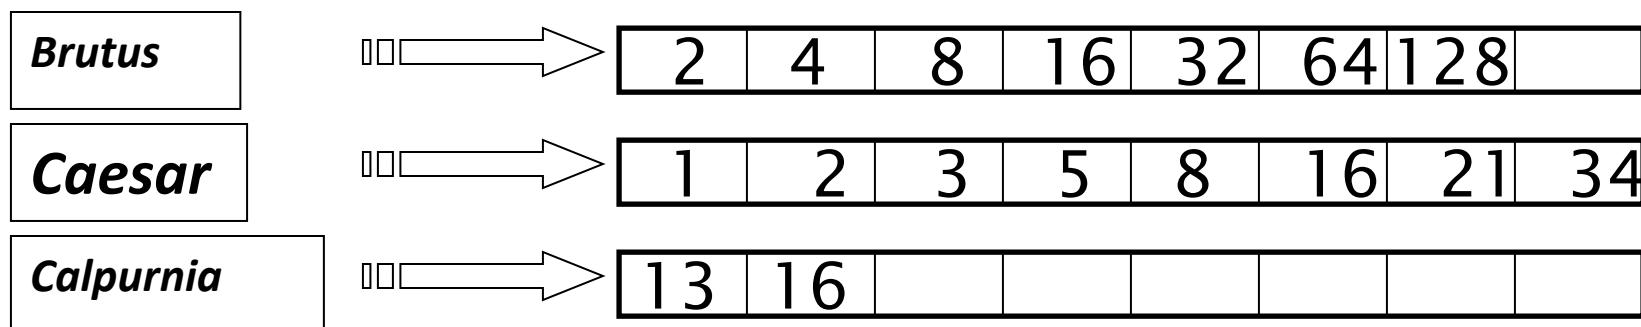
- The **Boolean retrieval model** is being able to ask a query that is a Boolean expression:
  - Boolean Queries are queries using *AND*, *OR* and *NOT* to join query terms
    - Views each document as a set of words
    - Is precise: document matches condition or not.
  - Perhaps the simplest model to build an IR system on
- Primary commercial retrieval tool for 3 decades.
- Many search systems you still use are Boolean:
  - Email, library catalog, Mac OS X Spotlight

# Example: WestLaw <http://www.westlaw.com/>

- Largest commercial (paying subscribers) legal search service (started 1975; ranking added 1992; new federated search (retrieves information from a variety of sources via a search application built on top of one or more search engines) added 2010)
- Tens of terabytes of data; ~700,000 users
- Majority of users *still* use boolean queries
- Example query:
  - What is the statute of limitations in cases involving the federal tort claims act?
  - **LIMIT! /3 STATUTE ACTION /S FEDERAL /2 TORT /3 CLAIM**
    - /3 = within 3 words, /S = in same sentence

# Query optimization

- What is the best order for query processing?
- Consider a query that is an *AND* of  $n$  terms.
- For each of the  $n$  terms, get its postings, then *AND* them together.



**Query:** *Brutus AND Calpurnia AND Caesar*

# Query optimization example

- Process in order of increasing freq:
  - start with smallest set, then keep cutting further.

This is why we kept  
document freq. in dictionary

Brutus	⇒	2   4   8   16   32   64   128
Caesar	⇒	1   2   3   5   8   16   21   34
Calpurnia	⇒	13   16

Execute the query as (*Calpurnia AND Brutus*) AND *Caesar*.

# More general optimization

- e.g., *(madding OR crowd) AND (ignoble OR strife) AND (killed OR slain)*
- Get doc. freq.'s for all terms.
- Estimate the size of each *OR* by the sum of its doc. freq.'s (conservative).
- Process in increasing order of *OR* sizes.

$\text{INTERSECT}(\langle t_1, \dots, t_n \rangle)$

```
1 terms  $\leftarrow$   $\text{SORTBYINCREASINGFREQUENCY}(\langle t_1, \dots, t_n \rangle)$ 
2 result  $\leftarrow$   $\text{postings}(\text{first}(\text{terms}))$ 
3 terms  $\leftarrow$   $\text{rest}(\text{terms})$ 
4 while terms  $\neq$  NIL and result  $\neq$  NIL
5   do result  $\leftarrow$   $\text{INTERSECT}(\text{result}, \text{postings}(\text{first}(\text{terms})))$ 
6   terms  $\leftarrow$   $\text{rest}(\text{terms})$ 
7 return result
```

# Exercise

- Recommend a query processing order for

*(tangerine OR trees) AND  
(marmalade OR skies) AND  
(kaleidoscope OR eyes)*

- Which two terms should we process first?

Term	Freq
eyes	213312
kaleidoscope	87009
marmalade	107913
skies	271658
tangerine	46653
trees	316812

# Does Google use the Boolean model?

- On Google, the default interpretation of a query [ $w_1 w_2 \dots w_n$ ] is  $w_1$  AND  $w_2$  AND ... AND  $w_n$
- Cases where you get hits that do not contain one of the  $w_i$  :
  - anchor text (clickable text)
  - page contains variant of  $w_i$  (morphology, spelling correction, synonym)
  - long queries ( $n$  large)
  - boolean expression generates very few hits
- Simple Boolean vs. Ranking of result set
  - Simple Boolean retrieval returns matching documents in no particular order.
  - Google (and most well designed Boolean engines) rank the result set – they rank good hits (according to some estimator of relevance) higher than bad hits.

# The extended Boolean model versus ranked retrieval

- *Ranked retrieval* - users largely use FREE TEXT QUERIES that is, just typing one or more words rather than using a precise language with operators for building up query expressions.
- Despite decades of academic research on the advantages of ranked retrieval, systems implementing the Boolean retrieval model were the main or only search option provided by large commercial information providers for three decades until the early 1990s
- However, these systems did not have just the basic Boolean operations (AND, OR, and NOT), but used extended Boolean retrieval models by incorporating additional operators
  - *proximity operator*
  - /s, /p, and /k ask for matches in the same sentence, same paragraph or within k words respectively.
  - Double quotes give a *phrase search* (consecutive words)
  - The exclamation mark (!) gives a trailing wildcard query, thus liab! matches all words starting with liab.
  - Additionally work-site matches any of *worksite*, *work-site* or *work site*

# Additional Features

1. We would like to better determine the set of terms in the dictionary and to provide retrieval that is **tolerant to spelling mistakes** and **inconsistent choice of words**.
2. It is often useful to search for **compounds or phrases** that denote a concept such as “**operating system**”. As the Westlaw examples show, we might also wish to do proximity queries such as **Gates NEAR Microsoft**. To answer such queries, the index has to be augmented to capture the proximities of terms in documents.
3. A Boolean model only records term presence or absence, but often we would like to accumulate evidence, giving more weight to documents that have a term several times as opposed to ones that contain it only once. To be able to do this we need *term frequency* information **TERM FREQUENCY** (the number of times a term occurs in a document) in postings lists.
4. Boolean queries just retrieve a set of matching documents, but commonly we wish to have an effective method **to order (or “rank”)** the returned results. This requires having a mechanism for determining a document score which encapsulates how good a match a document is for a query.