

Introduction to  
**Information Retrieval**

Lecture 9: Index Compression

# This lecture

---

BRUTUS	→	1	2	4	11	31	45	173	174
CAESAR	→	1	2	4	5	6	16	57	132
CALPURNIA	→	2	31	54	101				

- Collection statistics in more detail
  - How big are the dictionary and postings likely to be, for a given text documents collection?
- Dictionary compression
- Postings compression

# Why compression (in general)?

---

- Use less disk space
  - Saves a little money
- Keep more stuff in memory
  - Increases speed due to caching of more data
- Increase speed of data transfer from disk to memory
  - [read compressed data | decompress] is faster than [read uncompressed data]
  - Premise: Decompression algorithms are fast
    - True of the decompression algorithms we use

# Why compression for inverted indexes?

---

- Dictionary
  - Make it small enough to keep in main memory
  - Make it so small that you can keep some postings lists in main memory too
- Postings file(s)
  - Reduce disk space needed
  - Decrease time needed to read postings lists from disk
  - Large search engines keep a significant part of the postings in memory (compression lets you keep more in memory)
- We will devise various IR-specific compression schemes

# Sample text collection: Reuters RCV1

■ symbol	statistic	value
■ N	documents	800,000
■ L	avg. # tokens per doc	200
■ M	terms (= word types)	~400,000
■	avg. # bytes per token (incl. spaces/punct.)	6
■	avg. # bytes per token (without spaces/punct.)	4.5
■	avg. # bytes per term	7.5
■	non-positional postings	100,000,000

# Observations

---

- Preprocessing greatly affects the size of dictionary and number of postings
  - Stemming, case folding, stop word removal
- Percentage reduction can be different based on properties of the collections
  - E.g., lemmatizer for French reduces dictionary size much more than Porter stemmer for English

# Index parameters vs. what we index

(details *IIR* Table 5.1, p.80)

size of	word types (terms)			non-positional postings			positional postings		
	dictionary			non-positional index			positional index		
	Size (K)	Δ%	cumul %	Size (K)	Δ %	cumul %	Size (K)	Δ %	cumul %
Unfiltered	484			109,971			197,879		
No numbers	474	-2	-2	100,680	-8	-8	179,158	-9	-9
Case folding	392	-17	-19	96,969	-3	-12	179,158	0	-9
30 stopwords	391	-0	-19	83,390	-14	-24	121,858	-31	-38
150 stopwords	391	-0	-19	67,002	-30	-39	94,517	-47	-52
stemming	322	-17	-33	63,812	-4	-42	94,517	0	-52

Exercise: give intuitions for all the ‘0’ entries. Why do some zero entries correspond to big deltas in other columns?

# Lossless vs. lossy compression

---

- Lossless compression: All information is preserved.
  - What we mostly do in IR.
- Lossy compression: Discard some information
  - Makes sense when the discarded information is unlikely to be ever used by the IR system
- Several of the preprocessing steps can be viewed as lossy compression: case folding, stop words, stemming, number elimination.

# Compression

---

- We will consider compression schemes
  - Dictionary compression
  - Postings list compression

# DICTIONARY COMPRESSION

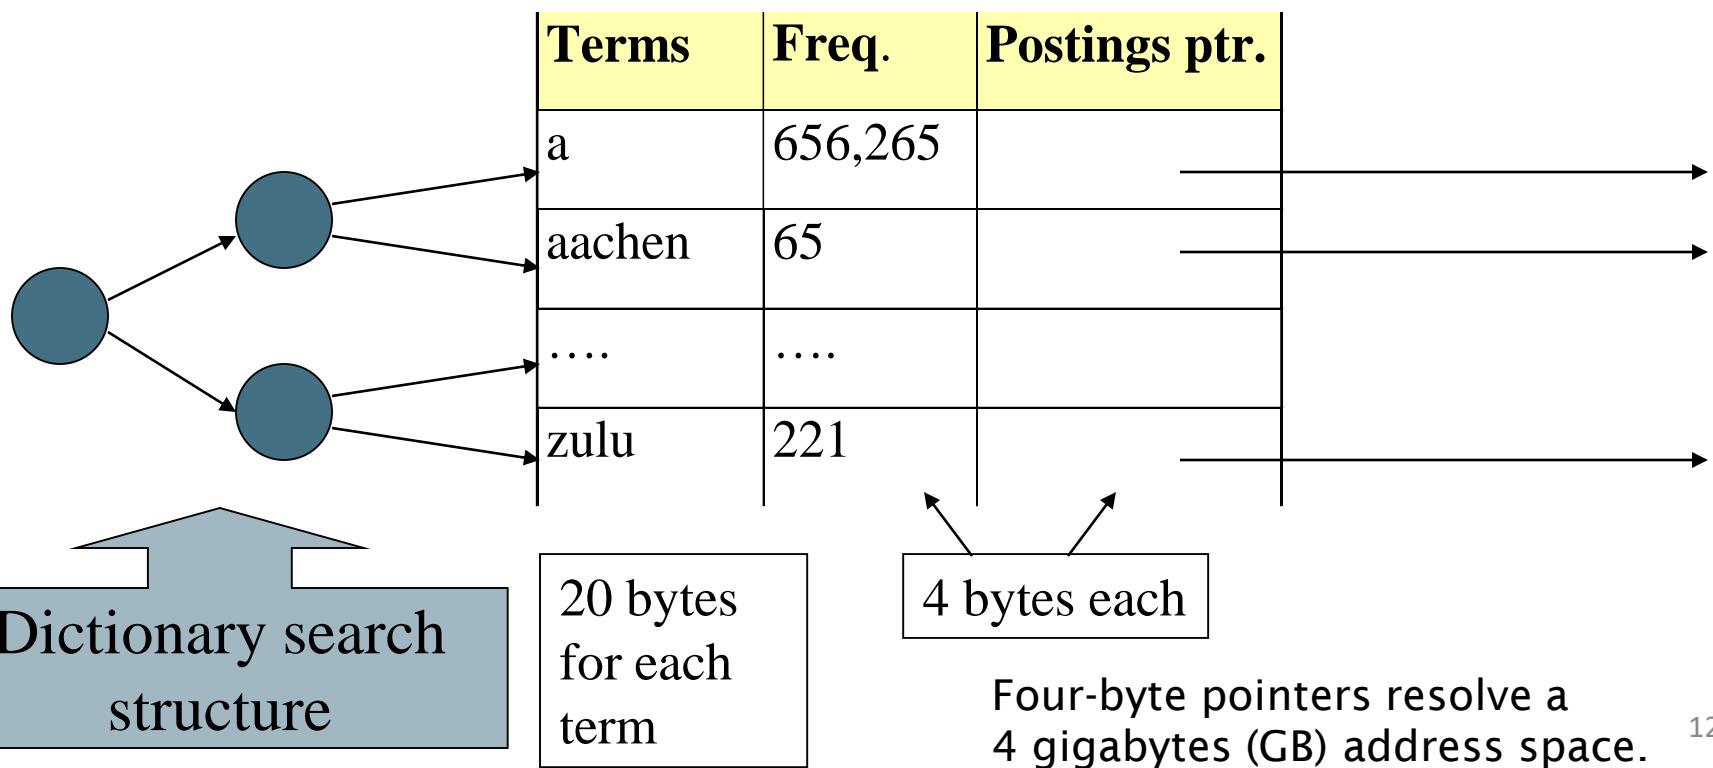
# Why compress the dictionary?

---

- Search begins with the dictionary
- We want to keep it in memory
- Memory footprint: competition with other applications
- Embedded/mobile devices may have very little memory
- Even if the dictionary isn't in memory, we want it to be small for a fast search startup time
- So, compressing the dictionary is important

# Dictionary storage

- Array of fixed-width entries
  - ~400,000 terms;  $28 \text{ bytes/term} = 11.2 \text{ MB}$ .



# Fixed-width terms are wasteful

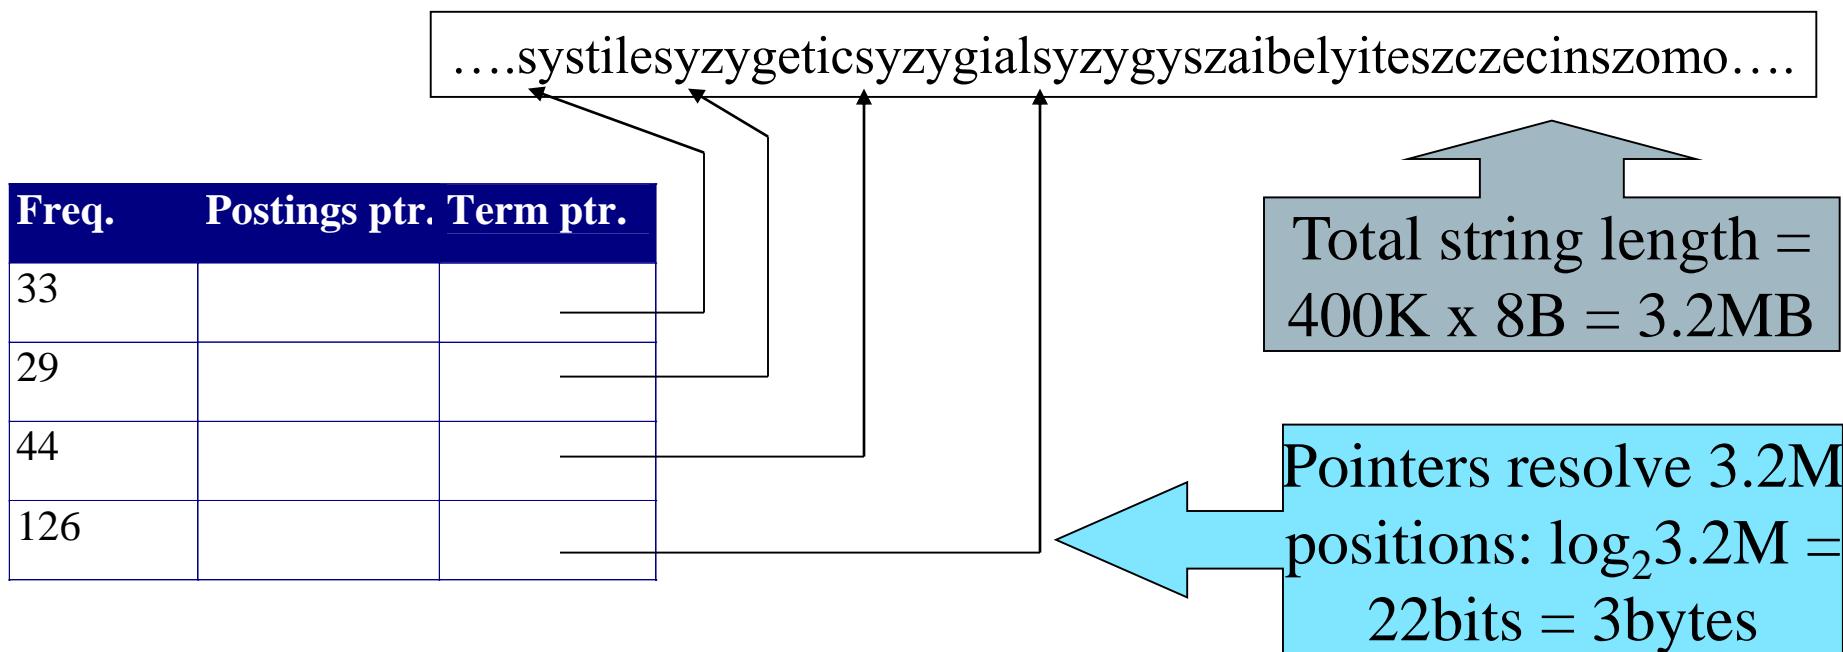
---

- Most of the bytes in the **Term** column are wasted – we allot 20 bytes even for 1 letter terms.
  - And we still can't handle *terms with more than 20 chars*
- Written English averages ~4.5 characters/word.
- Ave. dictionary word in English: ~8 characters
  - How do we use ~8 characters per dictionary term?

# Compressing the term list:

## Approach 1: Dictionary-as-a-String

- Store dictionary as a (long) string of characters:
  - Pointer to next word shows end of current word
  - Hope to save up to 60% of dictionary space.

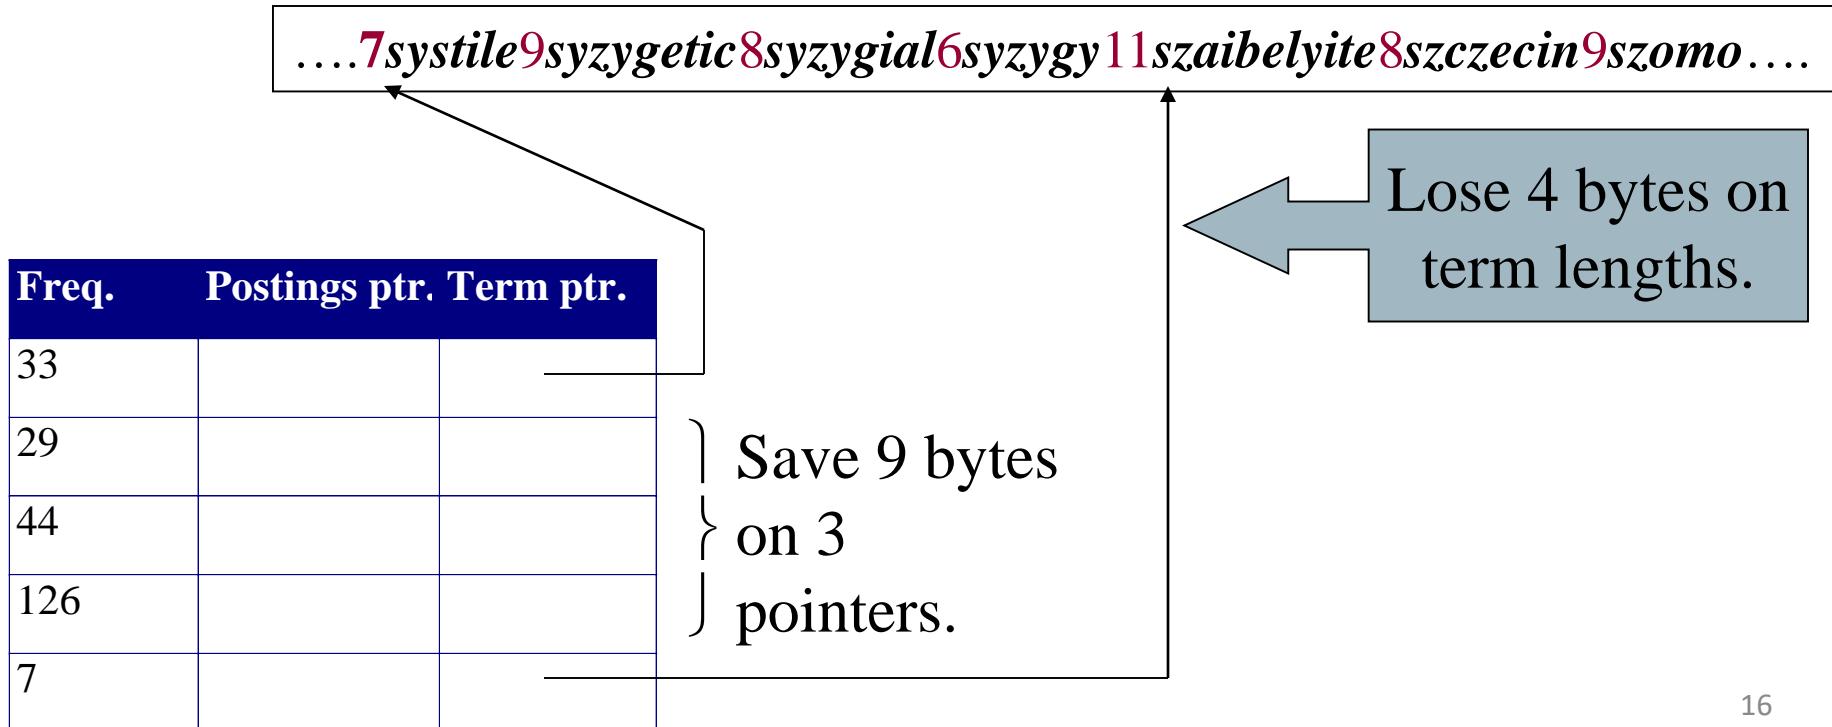


# Space for dictionary as a string

- 4 bytes per term for Freq.
  - 4 bytes per term for pointer to Postings.
  - 3 bytes per term pointer
  - Avg. 8 bytes per term in term string
  - 400K terms x 19  $\Rightarrow$  7.6 MB (against 11.2MB for fixed width)
- } Now avg. 11 bytes/term, not 20.

## Approach 2: Blocking

- Store pointers to every  $k$ -th term string.
  - Example below:  $k=4$ .
- Need to store term lengths (1 extra byte)



# Blocking

---

- Group terms into blocks, each having  $k$  terms
- Store a term pointer only for first term of each block
- Store the length of each term as one additional byte at the beginning of each term
- Search for terms in the compressed dictionary
  - Locate the term's block by **binary search**
  - Then locate term's position within the block by **linear search within the block**
- By increasing block size  $k$ : tradeoff between better compression and speed of term lookup

# Net saving

---

- Example for block size  $k = 4$
- Where we used 3 bytes/pointer without blocking
  - $3 \times 4 = 12$  bytes,

now we use  $3 + 4 = 7$  bytes.

Saved another  $\sim 0.5\text{MB}$  ( $400000 * 5/4$ ) . This reduces the size of the dictionary from 7.6 MB to 7.1 MB.

We can save more with larger  $k$ .

Why not go with larger  $k$ ?

By increasing the block size  $k$ , we get better compression. However, there is a tradeoff between compression and the speed of term lookup. By increasing  $k$ , we can get the size of the compressed dictionary arbitrarily close to the minimum of  $400,000 \times (4 + 4 + 1 + 8) = 6.8$  MB, but term lookup becomes prohibitively slow for large values of  $k$ .

## Approach 3: Front coding

- Front-coding:

- Sorted words commonly have long common prefix – store differences only
- In the case of Reuters, front coding saves another 1.2 MB,
- (for last  $k-1$  in a block of  $k$ )

**8automata8automate9automatic10automation**

$\rightarrow 8automat^*a1\diamond e2\diamond ic3\diamond ion$

Encodes *automat*

Extra length  
beyond *automat*.

Begins to resemble general string compression.

# RCV1: Our collection for this lecture

---

- As an example for applying scalable index construction algorithms, we will use the Reuters RCV1 collection.
  - This is one year of Reuters newswire (part of 1995 and 1996)
- The collection isn't really large enough, but it's publicly available and is a plausible example.

# A Reuters RCV1 document



You are here: Home > News > Science > Article

Go to a Section: U.S. International Business Markets Politics Entertainment Technology Sports Oddly Enough

## Extreme conditions create rare Antarctic clouds

Tue Aug 1, 2006 3:20am ET

[Email This Article](#) | [Print This Article](#) | [Reprints](#)



SYDNEY (Reuters) - Rare, mother-of-pearl colored clouds caused by extreme weather conditions above Antarctica are a possible indication of global warming, Australian scientists said on Tuesday.

[+] Text [-]

Known as nacreous clouds, the spectacular formations showing delicate wisps of colors were photographed in the sky over an Australian meteorological base at Mawson Station on July 25.

# Reuters RCV1 statistics

symbol	statistic	value
N	documents	800,000
L	avg. # tokens per doc	200
M	terms (= word types)	400,000
	avg. # bytes per token (incl. spaces/punct.)	6
	avg. # bytes per token (without spaces/punct.)	4.5
	avg. # bytes per term	7.5
	non-positional postings	100,000,000

4.5 bytes per word token vs. 7.5 bytes per word type: why?

# RCV1 dictionary compression summary

Technique	Size in MB
Fixed width	11.2
Dictionary-as-String with pointers to every term	7.6
Also, blocking $k = 4$	7.1
Also, Blocking + front coding	5.9

# **POSTINGS COMPRESSION**

# Postings compression

---

- The postings file is much larger than the dictionary, factor of at least 10.
- Key requirement: store each posting compactly.
- A posting for our purposes is a docID.
- For Reuters (800,000 documents), we would use 32 bits per docID when using 4-byte integers.
- Alternatively, we can use  $\log_2 800,000 \approx 20$  bits per docID.
- Our goal: use fewer than 20 bits per docID.

# Postings: two conflicting forces

---

- A term like *arachnocentric* occurs in maybe one doc out of a million – we would like to store this posting using  $\log_2 1M \sim 20$  bits.
- A term like *the* occurs in virtually every doc, so 20 bits/posting is too expensive.

# Postings file entry

- We store the list of docs containing a term in increasing order of docID.
  - *computer*: 33,47,154,159,202 ...
- Consequence: it suffices to *store gaps*.
  - 33,14,107,5,43 ...
- Hope: most gaps can be encoded/stored with far fewer than 20 bits.

# Three postings entries

	encoding	postings list							
THE	docIDs	...	283042	283043	283044	283045	...		
	gaps		1	1	1	1	...		
COMPUTER	docIDs	...	283047	283154	283159	283202	...		
	gaps		107	5	43	43	...		
ARACHNOCENTRIC	docIDs	252000	500100						
	gaps	252000	248100						

# Approach 1 : Variable length encoding

---

- Aim:
  - For *arachnocentric*, we will use  $\sim 20$  bits/gap entry.
  - For *the*, we will use  $\sim 1$  bit/gap entry.
- If the average gap for a term is  $G$ , we want to use  $\sim \log_2 G$  bits/gap entry.
- Key challenge: encode every integer (gap) with about as few bits as needed for that integer.
- This requires a *variable length encoding*
- Variable length codes achieve this by using short codes for small numbers

# Variable Byte (VB) codes

---

- For a gap value  $G$ , we want to use close to the fewest bytes needed to hold  $\log_2 G$  bits
- Begin with one byte to store  $G$  and **dedicate 1 bit in it to be a continuation bit  $c$**
- If  $G \leq 127$ , binary-encode it in the 7 available bits and set  $c = 1$
- Else encode  $G$ 's lower-order 7 bits and then use additional bytes to encode the higher order bits using the same algorithm
- At the end **set the continuation bit of the last byte to 1 ( $c = 1$ ) – and for the other bytes  $c = 0$ .**

# Example

docIDs	824	829	215406
gaps		5	214577
VB code	00000110 10111000	10000101	00001101 00001100 10110001

Postings stored as the byte concatenation

0000011010111000100001010000110010110001



Key property: VB-encoded postings are uniquely prefix-decodable.

Binary Rep. of 824

1100111000

Binary Rep. of 214577

110100011000110001

For a small gap (5), VB uses a whole byte.

# Other variable unit codes

---

- Instead of bytes, we can also use a different “unit of alignment”: 32 bits (words), 16 bits, 4 bits (nibbles).
- Variable byte alignment wastes space if you have many small gaps – nibbles do better in such cases.
- Variable byte codes:
  - Used by many commercial/research systems

# Variable bit-level codes: Unary code

- Represent  $n$  as  $n$  1s with a final 0.
  - Unary code for 3 is 1110.
  - Unary code for 40 is

- Unary code for 80 is:

- This doesn't look promising, but....

# Gamma codes

---

- We can compress better with bit-level codes
  - The Gamma code is the best known of these.
- Represent a gap  $G$  as a pair *length* and *offset*
- *offset* is  $G$  in binary, with the leading bit cut off
  - For example  $13 \rightarrow 1101 \rightarrow 101$
- *length* is the length of offset
  - For  $13$  (offset  $101$ ), this is  $3$ .
- We encode *length* with *unary code*:  $1110$ .
- Gamma code of  $13$  is the concatenation of *length* and *offset*:  $1110101$

# Gamma code examples

number	length	offset	$\gamma$ -code
0			none
1	0		0
2	10	0	10,0
3	10	1	10,1
4	110	00	110,00
9	1110	001	1110,001
13	1110	101	1110,101
24	11110	1000	11110,1000
511	11111110	11111111	11111110,11111111
1025	1111111110	0000000001	1111111110,0000000001

# Gamma code properties

---

- $G$  is encoded using  $2 \lfloor \log G \rfloor + 1$  bits
  - Length of offset is  $\lfloor \log G \rfloor$  bits
  - Length of length is  $\lfloor \log G \rfloor + 1$  bits
  - Eg.  $\lfloor \log 24 \rfloor = \lfloor 4.58 \rfloor = 4$
- All gamma codes have an odd number of bits
- Almost within a factor of 2 of best possible,  $\log_2 G$
- Gamma code is uniquely prefix-decodable, like VB
- Gamma code can be used for any distribution
- Gamma code is parameter-free

# RCV1 compression

Data structure	Size in MB
dictionary, fixed-width	11.2
dictionary, term pointers into string	7.6
with blocking, k = 4	7.1
with blocking & front coding	5.9
collection (text, xml markup etc)	3,600.0
collection (text)	960.0
Term-doc incidence matrix	40,000.0
postings, uncompressed (32-bit words)	400.0
postings, uncompressed (20 bits)	250.0
postings, variable byte encoded	116.0
postings, $\gamma$ -encoded	101.0

# Index compression summary

---

- We can now create an index for highly efficient Boolean retrieval that is very space efficient
- Only 4% of the total size of the collection