

Introduction to **Information Retrieval**

Lecture 5: Scoring, Term Weighting and the
Vector Space Model

Agenda

- Ranked retrieval
- Scoring documents
- Term frequency
- Collection statistics
- Weighting schemes
- Vector space scoring

Ranked retrieval

- Thus far, our queries have all been Boolean.
 - Documents either match or don't.
- Good for expert users with precise understanding of their needs and the collection
- Not good for the majority of users.
 - Most users incapable of writing Boolean queries (or they are, but they think it's too much work).
 - Most users don't want to wade through 1000s of results.
 - This is particularly true of web search.

Problem with Boolean search: feast or famine

- Boolean queries often result in either too few (=0) or too many (1000s) results.
- Query 1: “*standard user dlink 650*” → 200,000 hits
- Query 2: “*standard user dlink 650 no card found*”: 0 hits
- It takes a lot of skill to come up with a query that produces a manageable number of hits.
 - AND gives too few; OR gives too many

Ranked retrieval models

- Rather than a set of documents satisfying a query expression, in **ranked retrieval**, the system returns an ordering over the (top) documents in the collection for a query
- **Free text queries**: Rather than a query language of operators and expressions, the user's query is just one or more words in a human language
- *Ranked retrieval has been associated with free text queries.*

Feast or famine: not a problem in ranked retrieval

- When a system produces a ranked result set, large result sets are not an issue
 - Indeed, the size of the result set is not an issue
 - We just show the top k (≈ 10) results
 - We don't overwhelm the user
- ***Premise: the ranking algorithm works***

Scoring as the basis of ranked retrieval

- We wish to return in order the documents most likely to be useful to the searcher
- *How can we rank-order the documents in the collection with respect to a query?*
- Assign a score – say in $[0, 1]$ – to each document
- This score measures how well document and query “match”.

Parametric and Zone indexes


Metadata, Fields, Zones

- Documents can have metadata and fields
 - E.g., title of document, author of document, date of creation
- Zones similar to fields, but can contain arbitrary text
 - E.g., abstract, introduction, ... of a research paper
- We can have an index for each field/zone
 - To support queries like “documents having *merchant* in the title and *william* in the author list”
 - Either separate index for each field/zone, or part of the same index


Bibliographic Search

Search category	Value
Author	Example: Widom, J or Garcia-Molina
Title	Also a part of the title possible
Date of publication	Example: 1997 or <1997 or >1997 limits the search to the documents appeared in, before and after 1997 respectively
Language	Language the document was written in English <input type="button" value="▼"/>
Project	ANY <input type="button" value="▼"/>
Type	ANY <input type="button" value="▼"/>
Subject group	ANY <input type="button" value="▼"/>
Sorted by	Date of publication <input type="button" value="▼"/>
<input type="button" value="Start bibliographic search"/>	
Find document via ID <input type="text" value=" "/>	

► **Figure 6.1** Parametric search. In this example we have a collection with fields allowing us to select publications by zones such as Author and fields such as Language.



► **Figure 6.2** Basic zone index ; zones are encoded as extensions of dictionary entries.



► **Figure 6.3** Zone index in which the zone is encoded in the postings rather than the dictionary.

Weighted zone scoring

- Given a Boolean query q and a document d
 - Compute a 'zone match score' in $[0,1]$ for each zone/field of d with q
 - Compute **linear combination of zone match scores**, where each zone assigned a weight (sum of weights equal to 1.0)
 - Sometimes called 'ranked Boolean retrieval'
- How to decide the weights?
 - Option 1: Specified by experts, e.g., match in "title" has higher significance than match in "body"
 - Option 2: Learn from training examples – application of Machine Learning

- Given a Boolean query q and a document d , weighted zone scoring assigns to the pair (q, d) a score in the interval $[0, 1]$, by computing a linear combination of zone scores, where each zone of the document contributes a Boolean value. More specifically, consider a set of documents each of which has ℓ zones. Let $g_1, \dots, g_\ell \in [0, 1]$ such that

$$\sum_{i=1}^{\ell} g_i = 1. \text{ For } 1 \leq i \leq \ell,$$

For $1 \leq i \leq \ell$, let s_i be the Boolean score denoting a match (or absence thereof) between q and the i th zone.

- For instance, the Boolean score from a zone could be 1 if all the query term(s) occur in that zone, and zero otherwise.
- Then, the weighted zone score is defined to be

$$\sum_{i=1}^{\ell} g_i s_i.$$

- Consider the query Shakespeare in a collection in which each document has three zones: *author*, *title* and *body*.
- The Boolean score function for a zone takes on the value 1 if the query term Shakespeare is present in the zone, and zero otherwise.
- Weighted zone scoring in such a collection would require three weights g_1 , g_2 and g_3 , respectively corresponding to the *author*, *title* and *body* zones.
- Suppose we set $g_1 = 0.2$, $g_2 = 0.3$ and $g_3 = 0.5$ (so that the three weights add up to 1);
- Thus if the term Shakespeare were to appear in the *title* and *body* zones but not the *author* zone of a document, the score of this document would be 0.8.

-
- **Example 6.1:** Consider the query *shakespeare* in a collection in which each document has three zones: *author*, *title* and *body*. The Boolean score function for a zone takes on the value 1 if the query term *shakespeare* is present in the zone, and zero otherwise. Weighted zone scoring in such a collection would require three weights g_1 , g_2 and g_3 , respectively corresponding to the *author*, *title* and *body* zones.
 - **In the example above with weights $g_1 = 0.2$, $g_2 = 0.31$ and $g_3 = 0.49$, what are all the distinct score values a document may get?**

-
- Weighted zone scores can be directly computed from inverted indexes.
 - The algorithm treats the case when the query q is a two term query consisting of query terms q_1 and q_2 , and the Boolean function is AND: 1 if both query terms are present in a zone and 0 otherwise.

Implement the computation of weighted zone scores

$\text{INTERSECT}(p_1, p_2)$

```

1  answer ← ⟨ ⟩
2  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
3  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
4      then ADD( $\text{answer}, \text{docID}(p_1)$ )
5           $p_1 \leftarrow \text{next}(p_1)$ 
6           $p_2 \leftarrow \text{next}(p_2)$ 
7      else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
8          then  $p_1 \leftarrow \text{next}(p_1)$ 
9      else  $p_2 \leftarrow \text{next}(p_2)$ 
10 return  $\text{answer}$ 
```

► Figure 1.6 Algorithm for the intersection of two postings lists p_1 and p_2 .

Algorithm for computing the weighted zone score from two postings lists.

$\text{ZONESCORE}(q_1, q_2)$

```

1  float scores[N] = [0]
2  constant g[ $\ell$ ]
3   $p_1 \leftarrow \text{postings}(q_1)$ 
4   $p_2 \leftarrow \text{postings}(q_2)$ 
5  // scores[] is an array with a score entry for each document, initialized to zero.
6  //  $p_1$  and  $p_2$  are initialized to point to the beginning of their respective postings.
7  // Assume g[] is initialized to the respective zone weights.
8  while  $p_1 \neq \text{NIL}$  and  $p_2 \neq \text{NIL}$ 
9  do if  $\text{docID}(p_1) = \text{docID}(p_2)$ 
10     then scores[ $\text{docID}(p_1)$ ] ← WEIGHTEDZONE( $p_1, p_2, g$ )
11          $p_1 \leftarrow \text{next}(p_1)$ 
12          $p_2 \leftarrow \text{next}(p_2)$ 
13     else if  $\text{docID}(p_1) < \text{docID}(p_2)$ 
14         then  $p_1 \leftarrow \text{next}(p_1)$ 
15     else  $p_2 \leftarrow \text{next}(p_2)$ 
16 return scores
```

Learning weights

- How do we determine the weights g_i for weighted zone scoring?
 - These weights could be specified by an expert (or, in principle, the user);
 - these weights are “learned” using training examples
- *Machine-learned relevance.*
 1. We are provided with a set of *training examples*, each of which is a tuple consisting of a query q and a document d , together with a relevance judgment for d on q . In the simplest form, each relevance judgments is either *Relevant* or *Non-relevant*.
 2. The weights g_i are then “learned” from these examples, in order that the learned scores approximate the relevance judgments in the training examples.

- We now consider a simple case of weighted zone scoring, where each document has a *title* zone and a *body* zone.

$$\text{score}(d, q) = g \cdot s_T(d, q) + (1 - g)s_B(d, q).$$

- Training set examples* - each of which is a triple of the form

$$\Phi_j = (d_j, q_j, r(d_j, q_j)).$$

- Human editor delivers a relevance judgment $r(d_j, q_j)$ that is either *Relevant* or *Non-relevant*.

Example	DocID	Query	s_T	s_B	Judgment
Φ_1	37	linux	1	1	Relevant
Φ_2	37	penguin	0	1	Non-relevant
Φ_3	238	system	0	1	Relevant
Φ_4	238	penguin	0	0	Non-relevant
Φ_5	1741	kernel	1	1	Relevant
Φ_6	2094	driver	0	1	Relevant
Φ_7	3191	driver	1	0	Non-relevant

► Figure 6.5 An illustration of training examples.

s_T	s_B	Score
0	0	0
0	1	$1 - g$
1	0	g
1	1	1

► Figure 6.6 The four possible combinations of s_T and s_B .

- For each training example Φ_j we have Boolean values $sT(d_j, q_j)$ and $sB(d_j, q_j)$ that we use to compute a score

$$\text{score}(d_j, q_j) = g \cdot sT(d_j, q_j) + (1 - g)sB(d_j, q_j).$$

- We now compare this computed score to the human relevance judgment for the same document-query pair (d_j, q_j)
- Suppose that we define the error of the scoring function with weight g as

$$\varepsilon(g, \Phi_j) = (r(d_j, q_j) - \text{score}(d_j, q_j))^2,$$

- where we have quantized the editorial relevance judgment $r(d_j, q_j)$ to 0 or 1.
- Then, the total error of a set of training examples is given by

$$\sum_j \varepsilon(g, \Phi_j).$$

- The problem of learning the constant g from the given training examples then reduces to picking the value of g that minimizes the total error

Learning weights (simple machine learning)

Assume only two zones *title* (T) and *body* (B) with zone weights g and $1 - g$, respectively.

Example	DocID d	Query	s_T	s_B	Judgment (human expert)	r (quantized judgment)
ϕ_1	37	linux	1	1	Relevant	1
ϕ_2	37	penguin	0	1	Non-relevant	0
ϕ_3	238	system	0	1	Relevant	1
ϕ_4	238	penguin	0	0	Non-relevant	0
ϕ_5	1741	kernel	1	1	Relevant	1
ϕ_6	2094	driver	0	1	Relevant	1
ϕ_7	3191	driver	1	0	Non-relevant	0

Training examples

s_T	s_B	Score
0	0	0
0	1	$1 - g$
1	0	g
1	1	1

Four possible combinations of s_T and s_B and the corresponding
 $\text{score}(d, q) = g * s_T(d, q) + (1 - g) * s_B(d, q)$

Learning weights (simple machine learning)

Squared error of the scoring function with weight g on example ϕ is

$$\varepsilon(g, \phi) = (r(d, q) - \text{score}(d, q))^2$$

Example	d	Query	s_T	s_B	Score	r	ε	ε assuming $g = 0.4$
ϕ_1	37	linux	1	1	1	1	0	0
ϕ_2	37	penguin	0	1	$1 - g$	0	$(1 - g)^2$	0.36
ϕ_3	238	system	0	1	$1 - g$	1	g^2	0.16
ϕ_4	238	penguin	0	0	0	0	0	0
ϕ_5	1741	kernel	1	1	1	1	0	0
ϕ_6	2094	driver	0	1	$1 - g$	1	g^2	0.16
ϕ_7	3191	driver	1	0	g	0	g^2	0.16
							<u>ε</u>	<u>0.84</u>
							<u>Tot_</u> ε	<u>Tot_</u> $\varepsilon = 0.84$

s_T	s_B	Score	Score assuming $g = 0.4$
0	0	0	0
0	1	$1 - g$	0.6
1	0	g	0.4
1	1	1	1

-
- **Exercise 6.5**
 - Apply Equation 6.6 to the sample training set in Figure 6.5 to estimate the best value of g for this sample.
 - **Exercise 6.6**
 - For the value of g estimated in Exercise 6.5, compute the weighted zone score for each (query, document) example. How do these scores relate to the relevance judgments in Figure 6.5 (quantized to 0/1)?

$$\sum_j \varepsilon(g, \Phi_j).$$

Equation 6.4

Let n_{01r} (respectively, n_{01n}) denote the number of training examples for which $s_T(d_j, q_j) = 0$ and $s_B(d_j, q_j) = 1$ and the editorial judgment is *Relevant* (respectively, *Non-relevant*). Then the contribution to the total error in Equation (6.4) from training examples for which $s_T(d_j, q_j) = 0$ and $s_B(d_j, q_j) = 1$

is

$$[1 - (1 - g)]^2 n_{01r} + [0 - (1 - g)]^2 n_{01n}.$$

Learning weights (simple machine learning)

- Total error of a set of training examples $\text{Tot_}\varepsilon = \sum_j \varepsilon(g, \phi_j) = \sum_j (r(d_j, q) - \text{score}(d_j, q))^2$
- Goal is to choose g to minimize the total error.
- Note: Our example has only two zones with weights g and $1 - g$, respectively!
Generally, there will be l zones with weights g_1, \dots, g_l . Same principles!

s_T	s_B	Score	r	No.	ε
0	0	0	0	n_{00n}	0
0	0	0	1	n_{00r}	1
0	1	$1 - g$	0	n_{01n}	$(1 - g)^2$
0	1	$1 - g$	1	n_{01r}	g^2
1	0	g	0	n_{10n}	g^2
1	0	g	1	n_{10r}	$(1 - g)^2$
1	1	1	0	n_{11n}	1
1	1	1	1	n_{11r}	0

Total error $\text{Tot_}\varepsilon :$
$$(n_{01r} + n_{10n})g^2 + (n_{10r} + n_{01n})(1 - g)^2 + n_{00r} + n_{11n}$$

Learning weights (simple machine learning)

- Want to minimize total error $\text{Tot_}\varepsilon = (n_{01r} + n_{10n})g^2 + (n_{10r} + n_{01n})(1 - g)^2 + n_{00r} + n_{11n}$
- Differentiating w.r.t. g : $d(\text{Tot_}\varepsilon)/dg$
 $= 2(n_{01r} + n_{10n})g - 2(n_{10r} + n_{01n})(1 - g)$
- Find minimum by solving:

$$\begin{aligned} 2(n_{01r} + n_{10n})g - 2(n_{10r} + n_{01n})(1 - g) &= 0 \\ \rightarrow (n_{10r} + n_{10n} + n_{01r} + n_{01n})g &= n_{10r} + n_{01n} \\ \rightarrow g = (n_{10r} + n_{01n}) / (n_{10r} + n_{10n} + n_{01r} + n_{01n}) & \\ \rightarrow g = (0 + 1) / (0+1+2+1) = \frac{1}{4} = 0.25 & \end{aligned}$$

Learning weights (simple machine learning)

Squared error of the scoring function with weight g on example ϕ is

$$\varepsilon(g, \phi) = (r(d, q) - \text{score}(d, q))^2$$

Example	d	Query	s_T	s_B	Score	r	ε	ε assuming $g = 0.25$
ϕ_1	37	linux	1	1	1	1	0	0
ϕ_2	37	penguin	0	1	$1 - g$	0	$(1 - g)^2$	0.5625
ϕ_3	238	system	0	1	$1 - g$	1	g^2	0.0625
ϕ_4	238	penguin	0	0	0	0	0	0
ϕ_5	1741	kernel	1	1	1	1	0	0
ϕ_6	2094	driver	0	1	$1 - g$	1	g^2	0.0625
ϕ_7	3191	driver	1	0	g	0	g^2	0.0625
							<u>Tot_ε</u>	<u>$\text{Tot}_\varepsilon = 0.75$</u>

s_T	s_B	Score	Score assuming $g = 0.25$
0	0	0	0
0	1	$1 - g$	0.75
1	0	g	0.25
1	1	1	1

Term Frequency and Weighing

- Scoring has hinged on whether or not a query term is present in a zone within a document.
- Next logical step: a document or zone that mentions a query term more often has more to do with that query and therefore should receive a higher score.
- We assign to each term in a document a *weight* for that term, that depends on the number of occurrences of the term in the document.
- We would like to compute a score between a query term t and a document d , based on the weight of t in d .
- The simplest approach is to assign the weight to be equal to the number of occurrences of term t in document d .
- This weighting scheme is referred to as TERM FREQUENCY and is denoted $\text{tf}_{t,d}$, with the subscripts denoting the term and the document in order.

Recall: Binary term-document incidence matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	1	1	0	0	0	1
Brutus	1	1	0	1	0	0
Caesar	1	1	0	1	1	1
Calpurnia	0	1	0	0	0	0
Cleopatra	1	0	0	0	0	0
mercy	1	0	1	1	1	1
worser	1	0	1	1	1	0

Each document is represented by a binary vector $\in \{0,1\}^{|V|}$

Term-document count matrices

- Consider the number of occurrences of a term in a document:
 - Each document is a **count vector** in \mathbb{N}^v : a column below

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

Bag of words model

- For a document d , the set of weights determined by the tf weights above may be viewed as a quantitative digest of that document.
- *Bag of words model*, the exact ordering of the terms in a document is ignored but the number of occurrences of each term is material.
- In *John is quicker than Mary* and *Mary is quicker than John* have the same scores
- This is called the bag of words model.
- It seems intuitive that two documents with similar bag of words representations are similar in content.
- But, in a sense, this is a step back: The positional index was able to distinguish these two documents.
- Are all words in a document equally important?

Term frequency tf

- The term frequency $tf_{t,d}$ of term t in document d is defined as the number of times that t occurs in d .
- We want to use tf when computing query-document match scores. But how?
- Raw term frequency- all terms are considered equally important when it comes to assessing relevancy on a query.
- Raw term frequency is not what we want:
 - A document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term.
 - But not 10 times more relevant.
- Relevance does not increase proportionally with term frequency.

Document frequency

- Rare terms are more informative than frequent terms
 - Recall stop words
- Consider a term in the query that is rare in the collection (e.g., *arachnocentric*)
- A document containing this term is very likely to be relevant to the query *arachnocentric*
- → We want a high weight for rare terms like *arachnocentric*.

Document frequency, continued

- Frequent terms are less informative than rare terms
- Consider a query term that is frequent in the collection (e.g., *high*, *increase*, *line*)
- A document containing such a term is more likely to be relevant than a document that doesn't
- But it's not a sure indicator of relevance.
- → For frequent terms, we want positive weights for words like *high*, *increase*, and *line*
- But lower weights than for rare terms.
- Therefore, there is need to attenuate the effect of terms that occur too often in the collection for relevance determination.
- We will use **document frequency (df)** to capture this.
- *Document frequency dft*, is defined to be the number of documents in the collection that contain a term t .

Collection vs. Document frequency

- The **collection frequency** of t is the number of occurrences of t in the collection, counting multiple occurrences.
- The idea would be to reduce the tf weight of a term by a factor that grows with its collection frequency.
- Example:

Word	Collection frequency	Document frequency
<i>insurance</i>	10440	3997
<i>try</i>	10422	8760

- Which word is a better search term (and should get a higher weight)?
- For the purpose of scoring it is better to use a document-level statistic

Inverse document frequency - idf

- How is the document frequency df of a term used to scale its weight?
- Denoting the total number of documents in a collection as N , we define the *inverse document frequency* (idf) of a term t as follows:

$$\text{idf}_t = \log_{10} (N/\text{df}_t)$$

- df_t is the document frequency of t : the number of documents that contain t
- We use $\log (N/\text{df}_t)$ instead of N/df_t to “dampen” the effect of idf.

idf example, suppose $N = 1$ million

term	df_t	idf_t
calpurnia	1	6
animal	100	4
sunday	1,000	3
fly	10,000	2
under	100,000	1
the	1,000,000	0

$$\text{idf}_t = \log_{10} (N/\text{df}_t)$$

There is one idf value for each term t in a collection.

term	df_t	idf_t
car	18,165	1.65
auto	6723	2.08
insurance	19,241	1.62
best	25,235	1.5

► Figure 6.8 Example of idf values. Here we give the idf's of terms with various frequencies in the Reuters collection of 806,791 documents.

Thus the idf of a rare term is high, whereas the idf of a frequent term is likely to be low.

Effect of idf on ranking

- Does idf have an effect on ranking for one-term queries, like
 - iPhone
- idf has no effect on ranking one term queries
 - idf affects the ranking of documents for queries with at least two terms
 - For the query **capricious person**, idf weighting makes occurrences of **capricious** count for much more in the final document ranking than occurrences of **person**.

tf-idf weighting

- The definitions of term frequency and inverse document frequency are combined to produce a composite weight for each term in each document.
- The tf-idf weight of a term is the product of its tf weight and its idf weight.

$$\text{tf-idf}_{t,d} = \text{tf}_{t,d} \times \text{idf}_t.$$

- Best known weighting scheme in information retrieval
 - Note: the “-” in tf-idf is a hyphen, not a minus sign!
 - Alternative names: tf.idf, tf x idf

In other words, $\text{tf-idf}_{t,d}$ assigns to term t a weight in document d that is

1. highest when t occurs many times within a small number of documents (thus lending high discriminating power to those documents);
2. lower when the term occurs fewer times in a document, or occurs in many documents (thus offering a less pronounced relevance signal);
3. lowest when the term occurs in virtually all documents.

Score for a document given a query

- *Overlap score measure* : The score of a document d is the sum, over all query terms, of the number of times each of the query terms occurs in d .
- We can refine this idea so that we add up not the number of occurrences of each query term t in d , but instead the tf-idf weight of each term in d .

$$\text{Score}(q, d) = \sum_{t \in q} \text{tf-idf}_{t,d}.$$

Binary → count → weight matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf-idf weights $\in \mathbb{R}^{|V|}$

Documents as vectors

- We may view each document as a *vector* with one component corresponding to each term in the dictionary, together with a weight for each component that is given by tf-idf $_{t,d}$
- So we have a $|V|$ -dimensional vector space
- Terms are axes of the space
- Documents are points or vectors in this space
- Very high-dimensional space: tens of millions of dimensions in case of a web search engine
- These are very sparse vectors - most entries are zero.

Consider the table of term frequencies for 3 documents denoted Doc1, Doc2, Doc3 in Figure 6.9. Compute the tf-idf weights for the terms car, auto, insurance best, for each document, using the idf values from Figure 6.8.

Solution :

term	df_t	idf_t
car	18,165	1.65
auto	6723	2.08
insurance	19,241	1.62
best	25,235	1.5

► **Figure 6.8** Example of idf values. Here we give the idf's of terms with various frequencies in the Reuters collection of 806,791 documents.

	Doc1	Doc2	Doc3
car	27	4	24
auto	3	33	0
insurance	0	33	29
best	14	0	17

► **Figure 6.9** Table of tf values for Exercise 6.10.

$$\text{tf-idf} = \text{tf} * \text{idf}$$


terms	Doc1	Doc2	Doc3
Car	44.55	6.6	39.6
Auto	6.24	68.64	0
Insurance	0	53.46	46.98
Best	21	0	25.5

Formalizing vector space proximity

- First cut: distance between two points
 - (= distance between the end points of the two vectors)
- Euclidean distance?
- Euclidean distance is a bad idea . . .
- . . . because Euclidean distance is large for vectors of different lengths.
- Two documents having similar content can have large Euclidean distance simply because one document is much longer than the other

Why distance is a bad idea

The Euclidean distance between \vec{q} and $\vec{d_2}$ is large even though the distribution of terms in the query q and the distribution of terms in the document d_2 are very similar.



Use angle instead of distance

- Thought experiment: take a document d and append it to itself. Call this document d' .
- “Semantically” d and d' have the same content
- The Euclidean distance between the two documents can be quite large
- The angle between the two documents is 0, corresponding to maximal similarity.


- Key idea: Rank documents according to angle with query.

From angles to cosines

- The following two notions are equivalent.
 - Rank documents in increasing order of the angle between query and document
 - Rank documents in decreasing order of cosine(query,document)
- Cosine is a monotonically decreasing function for the interval $[0^\circ, 180^\circ]$

Cosine is a monotonically decreasing function

From angles to cosines



- But how – *and why* – should we be computing cosines?

Cosine is a monotonically decreasing function

cosine(query,document)

Dot product

Unit vectors

$$\cos(\vec{q}, \vec{d}) = \frac{\vec{q} \bullet \vec{d}}{\|\vec{q}\| \|\vec{d}\|} = \frac{\vec{q}}{\|\vec{q}\|} \bullet \frac{\vec{d}}{\|\vec{d}\|} = \frac{\sum_{i=1}^{|V|} q_i d_i}{\sqrt{\sum_{i=1}^{|V|} q_i^2} \sqrt{\sum_{i=1}^{|V|} d_i^2}}$$

q_i is the tf-idf weight of term i in the query
 d_i is the tf-idf weight of term i in the document

$\cos(\vec{q}, \vec{d})$ is the cosine similarity of \vec{q} and \vec{d} ... or,
equivalently, the cosine of the angle between \vec{q} and \vec{d} .

Cosine for length-normalized vectors


- For length-normalized vectors, cosine similarity is simply the dot product (or scalar product):

$$\cos(\vec{q}, \vec{d}) = \vec{q} \bullet \vec{d} = \sum_{i=1}^{|V|} q_i d_i$$

for q , d length-normalized.

This is helpful when the direction of the vector is meaningful but the magnitude is not

Cosine similarity illustrated



Computing vector scores

COSINESCORE(q)

- 1 float Scores[N] = 0
- 2 Initialize Length[N]
- 3 for each query term t
- 4 do calculate $w_{t,q}$ and fetch postings list for t
- 5 for each pair(d , $tf_{t,d}$) in postings list
- 6 do $Scores[d] += wf_{t,d} \times w_{t,q}$
- 7 Read the array Length[d]
- 8 for each d
- 9 do $Scores[d] = Scores[d] / Length[d]$
- 10 return Top K components of Scores[]

Variant tf-idf functions

- For assigning a weight for each term in each document, a number of alternatives to tf and tf-idf have been considered.

$$wf_{t,d} = \begin{cases} 1 + \log tf_{t,d} & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}.$$

In this form, we may replace tf by some other function wf as in obtain:

$$wf\text{-}idf_{t,d} = wf_{t,d} \times idf_t.$$

tf-idf example: Inc.ltc

Document: *car insurance auto insurance*

Query: *best car insurance*

Term	Query							Document				Prod
	tf-raw	tf-wt	df	idf	wt	n'liz e	tf-raw	tf-wt	wt	n'liz e	wt	
auto	0	0	5000	2.3	0	0	1	1	1	0.52	0	0
best	1	1	50000	1.3	1.3	0.34	0	0	0	0	0	0
car	1	1	10000	2.0	2.0	0.52	1	1	1	0.52	0.27	0.27
insurance	1	1	1000	3.0	3.0	0.78	2	1.3	1.3	0.68	0.53	0.53

Exercise: what is N , the number of docs?

$$\text{Doc length} = \sqrt{1^2 + 0^2 + 1^2 + 1.3^2} \gg 1.92$$

$$\text{Score} = 0+0+0.27+0.53 = 0.8$$

Maximum tf normalization

- We observe higher term frequencies in longer documents, merely because longer documents tend to repeat the same words over and over again.
- One well-studied technique is to normalize the tf weights of all terms occurring in a document by the maximum tf in that document.
- Normalized term frequency for each term t in document d is given by

$$\text{ntf}_{t,d} = a + (1 - a) \frac{\text{tf}_{t,d}}{\text{tf}_{\max}(d)},$$

where a is a value between 0 and 1 and is generally set to 0.4, although some early work used the value 0.5. The term a is a *smoothing* term.

- Suppose we were to take a document d and create a new document d' by simply appending a copy of d to itself. While d' should be no more relevant to any query than d is, the use of

$$\text{Score}(q, d) = \sum_{t \in q} \text{tf-idf}_{t,d}.$$

would assign it twice as high a score as d . Replacing $\text{tf-idf}_{t,d}$ by $\text{ntf-idf}_{t,d}$ eliminates the anomaly in this example.

-
- Maximum tf normalization does suffer from the following issues:
 - a change in the stop word list can dramatically alter term weightings (and therefore ranking).
 - a document may contain an unusually large number of occurrences of a term, not representative of the content of that document
 - a document in which the most frequent term appears roughly as often as many other terms should be treated differently from one with a more skewed distribution.

Document and query weighting schemes

tf-idf weighting has many variants

Term frequency	Document frequency	Normalization
n (natural) $tf_{t,d}$	n (no) 1	n (none) 1
I (logarithm) $1 + \log(tf_{t,d})$	t (idf) $\log \frac{N}{df_t}$	c (cosine) $\frac{1}{\sqrt{w_1^2 + w_2^2 + \dots + w_M^2}}$
a (augmented) $0.5 + \frac{0.5 \times tf_{t,d}}{\max_t(tf_{t,d})}$	p (prob idf) $\max\{0, \log \frac{N - df_t}{df_t}\}$	u (pivoted unique) $1/u$
b (boolean) $\begin{cases} 1 & \text{if } tf_{t,d} > 0 \\ 0 & \text{otherwise} \end{cases}$		b (byte size) $1/CharLength^\alpha, \alpha < 1$
L (log ave) $\frac{1 + \log(tf_{t,d})}{1 + \log(\text{ave}_{t \in d}(tf_{t,d}))}$		

Columns headed ‘n’ are acronyms for weight schemes.

Weighting may differ in queries vs documents

- Many search engines allow for different weightings for queries vs. documents
- SMART Notation: denotes the combination in use in an engine, with the notation *ddd.ooo*, using the acronyms from the previous table
- A very standard weighting scheme is: Inc.ltc
- Document: logarithmic tf (l as first character), no idf and cosine normalization
- Query: logarithmic tf (l in leftmost column), idf (t in second column), cosine normalization ...

Pivoted normalized document length

- We normalized each document vector by the Euclidean length of the vector, so that all document vectors turned into unit vectors.
- In doing so, we eliminated all information on the length of the original document; thus masking some subtleties about longer documents.
- We introduce a form of normalizing the vector representations of documents in the collection, so that the resulting “normalized” documents are not necessarily of unit length.
- Then, when we compute the dot product score between a (unit) query vector and such a normalized document, the score is skewed to account for the effect of document length on relevance.
- This form of compensation for document length is known as *pivoted document length normalization*.

- Suppose that we were given, for each query q and for each document d , a Boolean judgment of whether or not d is relevant to the query q ; we may compute a *probability of relevance* as a function of document length, averaged over all queries in the ensemble. The resulting plot may look like the curve drawn in thick lines
 - To compute this curve, we bucket documents by length and compute the fraction of relevant documents in each bucket, then plot this fraction against the median document length of each bucket.
 - On the other hand, the curve in thin lines shows the same documents and query ensemble if we were to use relevance as prescribed by cosine normalization – thus, cosine normalization
- The thin and thick curves crossover at a point p corresponding to document length ℓ_p , which we refer to as the *pivot length*;


- Next “rotate” the cosine normalization curve counter-clockwise about p so that it more closely matches thick line representing the relevance vs. document length curve.
 - we do so by using in Equation

$$\text{score}(q, d) = \frac{\vec{V}(q) \cdot \vec{V}(d)}{|\vec{V}(q)| |\vec{V}(d)|}.$$


- a normalization factor for each document vector $\vec{V}(d)$ that is not the Euclidean length of that vector, but instead one that is larger than the Euclidean length for documents of length less than lp , and smaller for longer documents.
- Pivoted length normalization

$$au_d + (1 - a)\text{piv},$$

- Where a is the slope and u_d is the number of unique terms in document d .



► Figure 6.16 Pivoted document length normalization.



► Figure 6.17 Implementing pivoted document length normalization by linear scaling.

Summary – vector space ranking

- Represent the query as a weighted tf-idf vector
- Represent each document as a weighted tf vector
- Compute the cosine similarity score for the query vector and each document vector
- Rank documents with respect to the query by score
- Return the top K (e.g., $K = 10$) to the user